

Host A

Host B

Seq=92, 8 bytes of data

Seq=100, 20 bytes of data

ACK=100

ACK=100

ACK=100

ACK=100

Seq=100, 20 bytes of data

fast retransmit after sender receipt of triple duplicate ACK

**Flow Control:** endpoints in TCP prevent a fast sender from “overwhelming” a slow receiver. Each w/**receive window** sizes.  
\* Receiver controls sender, so sender won't overflow receiver's buffer by transmitting too much too fast.

**Receive Window (rwnd):**  
+ two hosts exchange rwnd size at beginning  
+ sender guarantees that diff between last seq # & ACK sent never > rwnd

**Principles of Congestion Control**  
+ “too many sources sending too much data too fast for the network to handle”  
+ manifestations:  
- lost pkts (buffer overflow at routers)  
- long delays (queuing in router buffers)

**Causes of Congestion (from ex in book):**  
+ too many senders, too many receivers  
+ one router, w/infinite buffers  
+ output link capacity: R  
+ no retransmission  
Perfect knowledge scenario:  
+ sender sends only when avail buffer  
Known Loss Scenario w/out free buff space:  
+ pkts can be lost, dropped @ router bc full buffers  
+ sender only resends if packet known to be lost  
Known Loss w/free buffer space:  
+ packets can be lost, dropped at router bc full buffer  
+ sender only resends if packet known to be lost  
Duplicates (realistic):  
+ packets can be lost, dropped at router bc full buffer  
+ sender times out permanently, sends 2 copies, delivers

**Costs of congestion:**  
+ more work (retrans) for given “goodput”  
+ unneeded retrans: link carries multiple copies of pkt\*  
- dec goodput \*  
+ when pkt dropped, any “upstream trans capacity used” \*  
for that pkt was wasted! \*


**TCP Congestion Control: add inc/mult dec** \*  
+ sender inc trans rate (window size), until loss occurs  
- additive inc: inc cwnd by 1 MSS every RTT until loss detected  
- multiplicative dec: cut cwnd in half after loss

+ **Sending Rate:** roughly: send cwnd bytes, wait RTT for ACKS, then send more bytes ; rate = cwnd/ RTT (bytes/sec)  
+ **Last byte sent/ACK <= cwnd**

\*\*\*\*\*

\* **Congestion control** is different from flow control. It uses observations about network behavior to control the amount of data that is sent. A variable called the “congestion window” (cwnd) is set to the maximum number of bytes that can be transmitted at one time. The congestion window grows when acknowledgements are received in a timely way and shrinks when timeouts occur.  
\* Seq, length, and ACK fields in the TCP segment. When the receiver acknowledges a segment with seq and len, it send the ACK=seq + len, the number of the next byte expected. (NOTE: this happens only if all segments up to seq have been received; if there are gaps, then the acknowledgement is for the last segment that was received with no preceding gaps.)

**\*rwnd Question (assume rwnd=0): How does recv tell sender to start sending again?**  
\*  
\* Answer: sender sends tiny messages (1 byte each); recv then uses  
\* ACK messages for these to carry info about rwnd.



**TCP Slow Start**  
+ when connection begins, inc rate  
expon until 1<sup>st</sup> loss event  
- init cwnd = 1 MSS  
- double cwnd every RTT  
- done by incrementing cwnd per ACK received  
- **TCP Tahoe, cwnd = 1**  
- **TCP Reno, 3 dupACKs, cwnd=cwnd/2**  
- Slow Start Treshold  
**ssthresh = cwnd/2**  
**\*3 dupAcks = ssthresh=cwnd/2 +3**

The rate of growth of cwnd is determined by another variable **ssthresh**, the slow-start threshold. When the value of cwnd is than ssthresh, the size of the window grows exponentially (doubling). When cwnd exceeds ssthresh, the window size grows only linearly.